

# Chapter 9: Game Tree Search



## Recap on previous chapter

---

- Stackelberg game: **Player 1: leader**, **Player 2: follower**
- Stackelberg vs strategy game
- Strategy game and stackelberg game in zero-sum
- Stackelberg for general 2-persons game (LP)
- In ( $>$ )3-player normal form games, an optimal Stackelberg strategy is an NP-Hard problem
- Stackelberg Competition
- Security Game

## Recap on previous chapter

---

		Follower					
		L		M		R	
Leader	U	1	1	3	0	2	2
	M	0	0	2	1	1	1
	D	3	4	2	3	3	0

Present a solution for the optimal Stackelberg strategy

# Finite perfect-information zero-sum games

---

- Finite: finitely players, actions, states
- Perfect information: every player knows the current state, all of the strategies, and what they do
- No simultaneous actions – players move one-at-a-time
- Constant-sum: regardless of how the game ends,  
 $\Sigma \text{ players' payoff} = k$ :
  - Zero-sum game when  $k = 0$
  - Thus constant-sum games usually are called zero-sum games

# Settings

---

- Two players whose actions alternate
- Utility values for each player are the opposite of the other
  - This creates the adversarial situation
- Fully observable environments
- Deterministic game: go, chess
- Generalizes to stochastic games: backgammon, monopoly, yahtzee, parcheesi, roulette, craps

# A brief history

---

1846: machine to play tic-tac-toe

1928: minimax theorem (von Neumann)

1944: backward-induction algorithm to produce perfect play (von Neumann & Morgenstern)

1950: minimax algorithm -finite horizon, approximate evaluation (Shannon)

1951: program (on paper) for playing chess (Turing)

1956: pruning to allow deeper search


1957: first complete chess program

1967: first program to compete in human chess tournaments

1997 (IBM): Deep Blue

2017: (Deepmind): AlphaGo

# Game tree (2-player, deterministic, turns)



Root node = the initial state

Non-terminal node = the states a player can move to

Terminal node = a state where the game ends

**How do we search this tree to find the optimal move?**

# Search versus Games

---

- Search – no adversary
  - Solution is (heuristic) method for finding goal
  - Heuristics and techniques can find *optimal* solution
  - Evaluation function: estimate of cost from start to goal through given node
- Games – adversary
  - Solution is strategy
    - strategy specifies move for every possible opponent reply
  - Time limits force an *approximate* solution
  - Evaluation function: evaluate “goodness” of game position
  - Examples: chess, checkers, go...



# Two players: MAX and MIN

---

MAX moves first and take turns until the game is over

- Winner gets reward, loser gets penalty.
- Zero sum means the sum of the reward and the penalty is a constant
- **Formal definition as a search problem:**
  - **Initial state:** Set-up specified by the rules, e.g., initial board of chess
  - **Player:** Defines which player has the move in a state
  - **Actions:** the set of legal moves in a state
  - **Result:** Transition model defines the result state of a move
  - **Successor function:** list of (move,state) pairs specifying legal moves
  - **Terminal-Test:** Is the game finished? True if finished, false other.
  - **Utility function:** Gives numerical value of terminal state for player
    - E.g., win (+1), lose (-1), and draw (0) in tic-tac-toe.
    - E.g., win (+1), lose (0), and draw (1/2) in chess.
- MAX uses search tree to determine next move.

# An optimal procedure: The Min-Max method


---

Find the optimal strategy for Max:

1. Generate the whole game tree, down to the leaves
2. Apply utility (payoff) function to each leaf
3. Back-up values from leaves through branch nodes:
  - a Max node computes the maximum of its child values
  - a Min node computes the minimum of its child values
4. At root: choose the move leading to the child of highest value

# Game Trees


---



A two-players game tree: the  $\triangle$  and  $\nabla$  nodes are moves by MAX and MIN, respectively. The terminal nodes show the payoff for MAX, whereas the payoffs of other nodes can be computed by the minimax algorithms

# Two-Players Game Tree

---




Backward induction

# Two-Players Game Tree


---

Maximizes the payoffs for MAX



# Two-Players Game Tree

---



# Pseudocode for Minimax Algorithm

---

**function** MINIMAX-DECISION(*state*) **returns** *an action*

**inputs:** current state in game

**return**  $\arg \max_{a \in \text{ACTIONS}(\textit{state})} \text{MIN-VALUE}(\text{Result}(\textit{state}, a))$

---

**function** MAX-VALUE(*state*) **returns** *a utility value*

**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

**for** *a* in ACTIONS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{Result}(\textit{state}, a)))$

**return** *v*

---

**function** MIN-VALUE(*state*) **returns** *a utility value*

**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow +\infty$

**for** *a* in ACTIONS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{Result}(\textit{state}, a)))$

**return** *v*

# Properties of Minimax Algorithm

---

- Complete?
  - Yes (if tree is finite).
- Optimal?
  - Yes (against an optimal opponent).
  - Can it be beaten by an opponent playing sub-optimally?
    - No. (Why not?)
- Time complexity?
  - $O(b^m)$  [**b =max. number of child., m =max. depth of any node**]
- Space complexity?
  - $O(bm)$  (depth-first search, generate all actions at once)



# Game Tree Size

---

- Tic-Tac-Toe
  - $b \approx 5$  legal actions per state on average, total of 9 depth in game
  - $5^9 = 1,953,125$
  - **exact solution quite reasonable**
- Chess
  - $b \approx 35$  (approximate average branching factor)
  - $d \approx 100$  (depth of game tree for “typical” game)
  - $b^d \approx 35^{100} \approx 10^{154}$  nodes!!
  - **exact solution completely infeasible**

**It is usually impossible to search the whole tree**


# (Static) Heuristic Evaluation Functions

---

- An Evaluation Function
  - Estimates how good the current state is for a player
  - Evaluate how good it is for the player, how good it is for the opponent, then subtract the opponent's score from the player's
  - Called “static” because it is called on a static board position
  - Othello: Number of white pieces - Number of black pieces
  - Chess: Value of all white pieces - Value of all black pieces
- Typical values from  $-\infty$  (loss) to  $+\infty$  (win) or  $[-1, +1]$
- If the board evaluation is  $X$  for a player, it's  $-X$  for the opponent e.g., zero-sum game


# Evaluation Functions (example)

---



Black to move

White slightly better



White to move

Black winning

For chess, typical linear weighted sum of features:

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \cdots + w_n f_n(s)$$

e.g.,  $w_1 = 9$  with  $f_1(s)$ =number of white queens - number of black queens, etc.

# Applying MiniMax to tic-tac-toe


- The static heuristic evaluation function: current state  $n$

$M(n)$ : the number of my possible winning lines

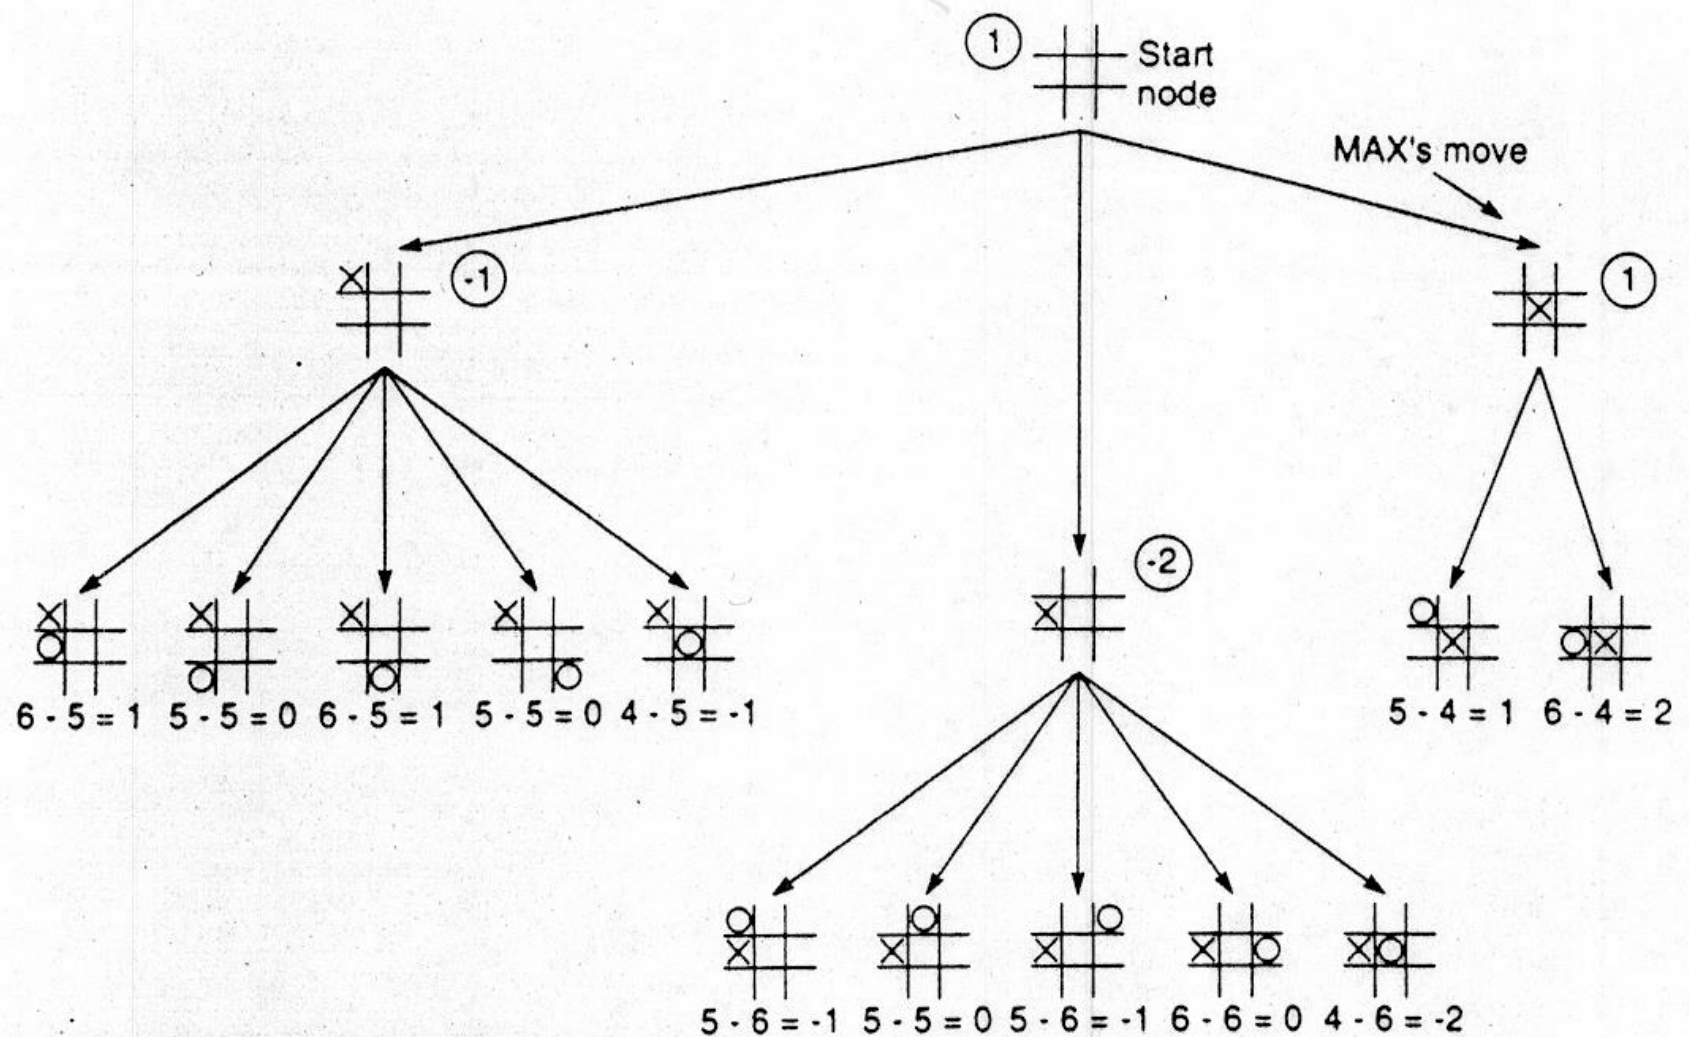
$O(n)$ : the number of the opponent's possible winning lines

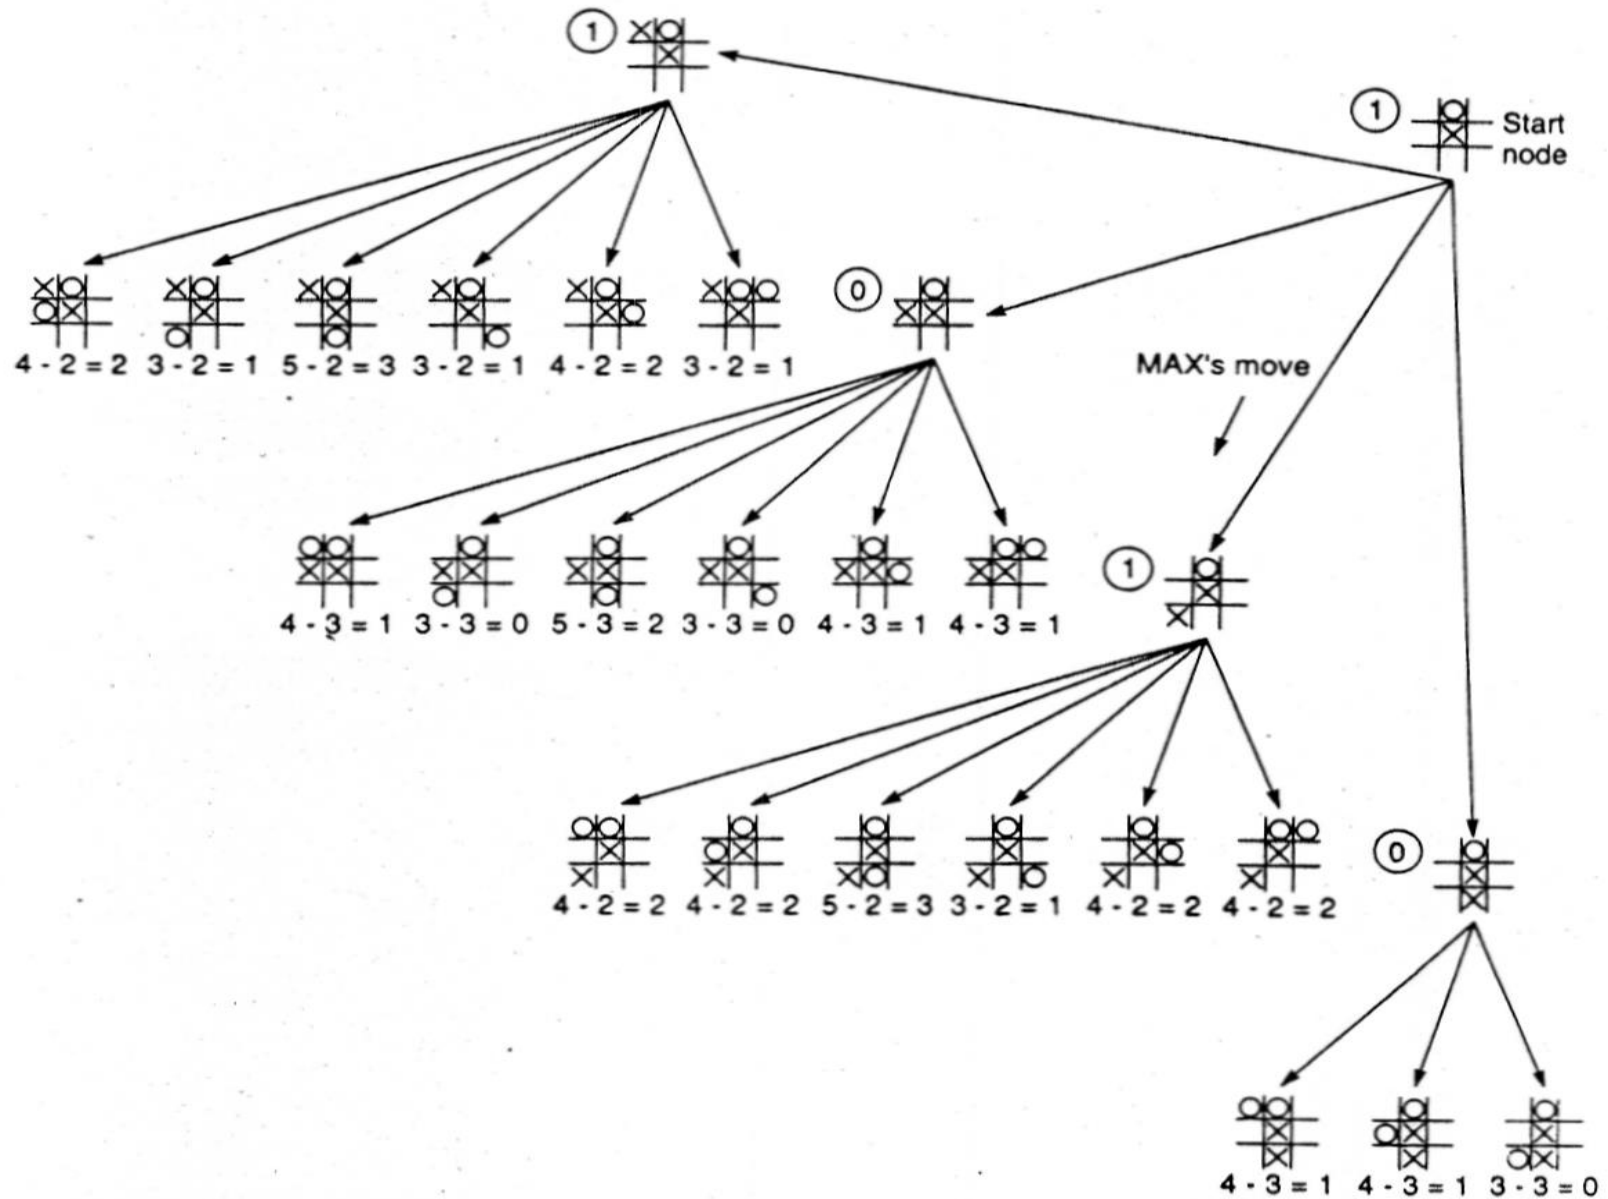
Heuristic evaluation function

$$E(n) = M(n) - O(n)$$

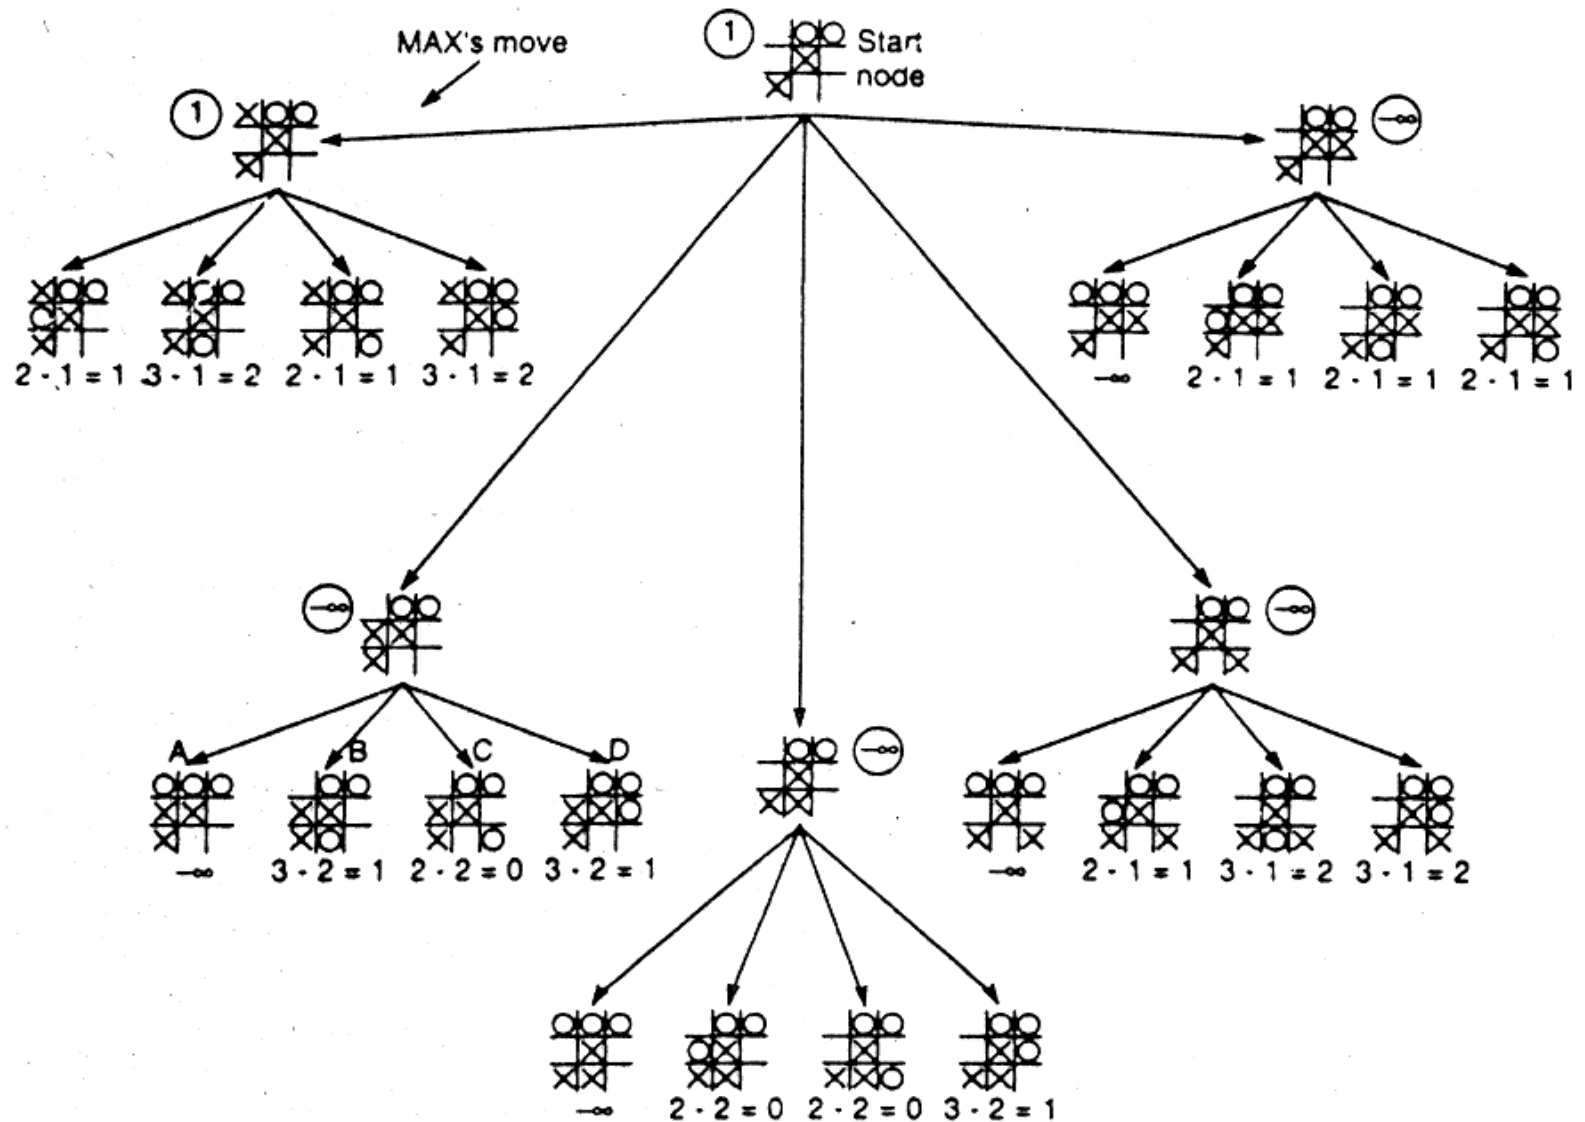


# Backup Values





# Near the end game



# Chapter 9: Game Tree Search





## Recap on last course

---

- Two players made adversary actions alternately
- Game problem  $\rightarrow$  search problem
- Minimax algorithm
- Evaluation function

# Applying MiniMax to tic-tac-toe


- The static heuristic evaluation function: current state  $n$

$M(n)$ : the number of my possible winning lines

$O(n)$ : the number of the opponent's possible winning lines

Heuristic evaluation function


$$E(n) = M(n) - O(n)$$



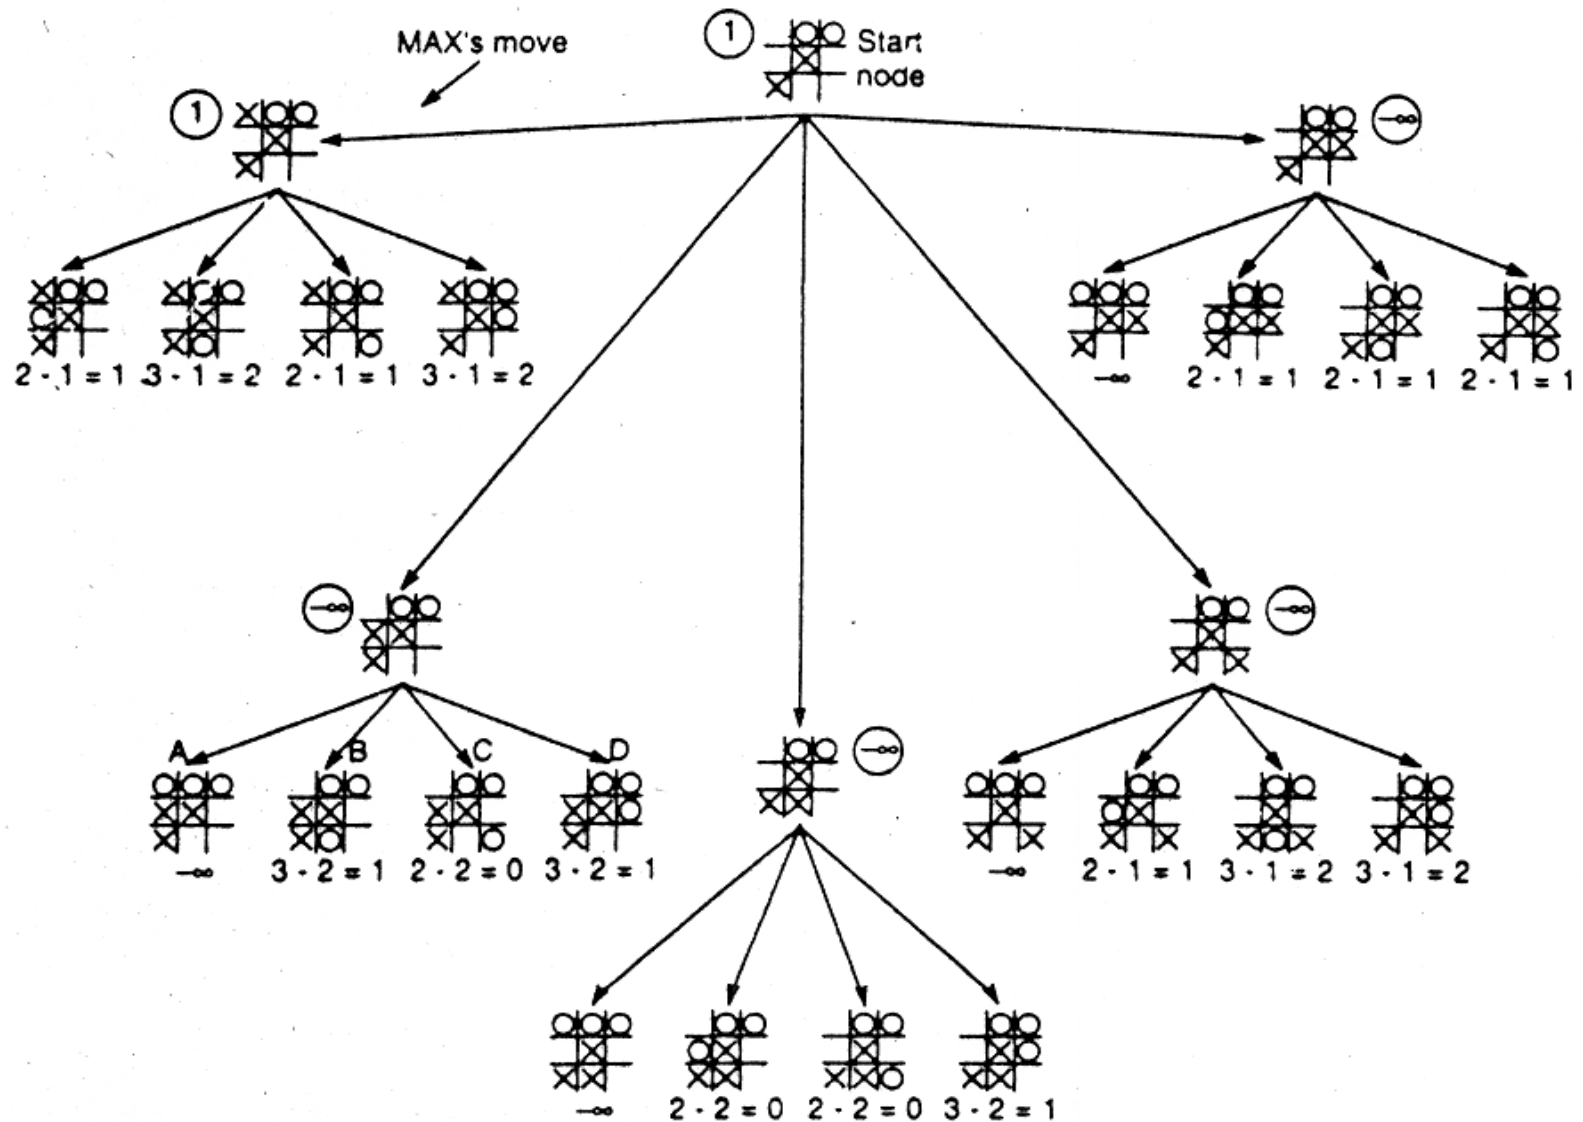
# Alpha-Beta Pruning: Exploiting the fact of an Adversary

- If a position is provably bad:
  - It is NO USE expending search time to find out exactly how bad
- If the adversary can force a bad position:
  - It is NO USE expending search time to find out the good positions that the adversary won't let you achieve anyway
- Bad = not better than we already know we can achieve elsewhere
- Contrast normal search:
  - ANY node might be a winner.
  - ALL nodes must be considered.

# Tic-Tac-Toe Example with Alpha-Beta Pruning




# Tic-Tac-Toe Example with Alpha-Beta Pruning



# Another Alpha-Beta Example


---

Do DF-search until first leaf




# Alpha-Beta Example (continued)

---



# Alpha-Beta Example (continued)


---






# Alpha-Beta Example (continued)

---




# Alpha-Beta Example (continued)

---




# Alpha-Beta Example (continued)

---




# Alpha-Beta Example (continued)

---



# Alpha-Beta Example (continued)

---



# Alpha-beta algorithm

---

- Depth first search
  - only considers nodes along a single path from root at any time
  - $\alpha$  = highest-value choice found at any choice point of path for MAX (initially,  $\alpha = -\infty$ )
  - $\beta$  = lowest-value choice found at any choice point of path for MIN (initially,  $\beta = +\infty$ )
- Pass current values of  $\alpha$  and  $\beta$  down to child nodes during search.
- Update values of  $\alpha$  and  $\beta$  during search:
  - MAX updates  $\alpha$  at MAX nodes
  - MIN updates  $\beta$  at MIN nodes
- Prune remaining branches at a node when  $\alpha \geq \beta$

# When to Prune

---

- Prune whenever  $\alpha \geq \beta$ 
  - Prune below a MAX node whose  $\alpha$  value becomes greater than or equal to the  $\beta$  value of its ancestors
    - **Max nodes update  $\alpha$**  based on children's returned values.
  - Prune below a Min node whose  $\beta$  value becomes less than or equal to the  $\alpha$  value of its ancestors
    - **Min nodes update  $\beta$**  based on children's returned values.

# Pseudocode for Alpha-Beta Algorithm

---

**function** ALPHA-BETA-SEARCH(*state*) **returns** *an action*

**inputs:** *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\textit{state}, -\infty, +\infty)$

**return** the *action* in  $\text{ACTIONS}(\textit{state})$  with value  $v$

---

**function**  $\text{MAX-VALUE}(\textit{state}, \alpha, \beta)$  **returns** *a utility value*

**if**  $\text{TERMINAL-TEST}(\textit{state})$  **then return**  $\text{UTILITY}(\textit{state})$

$v \leftarrow -\infty$

**for**  $a$  in  $\text{ACTIONS}(\textit{state})$  **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{Result}(s, a), \alpha, \beta))$

**if**  $v \geq \beta$  **then return**  $v$

$\alpha \leftarrow \text{MAX}(\alpha, v)$

**return**  $v$


---

(MIN-VALUE is defined analogously)




# Alpha-Beta Example Revisited

Do DF-search until first leaf




# Alpha-Beta Example (continued)

---




# Alpha-Beta Example (continued)

---




# Alpha-Beta Example (continued)

---




# Alpha-Beta Example (continued)

---




# Alpha-Beta Example (continued)

---




# Alpha-Beta Example (continued)

---




# Alpha-Beta Example (continued)






# Alpha-Beta Example (continued)

---




# Alpha-Beta Example (continued)

---




# Alpha-Beta Example (continued)

---




# Alpha-Beta Example (continued)

---



# Alpha-Beta Example (continued)

---



# Effectiveness of Alpha-Beta Search


---

- Worst Case
  - branches are ordered so that no pruning takes place. In this case alpha-beta gives no improvement over exhaustive search
- Best Case
  - each player's best move is the left-most child (i.e., evaluated first)
  - in practice, performance is closer to best rather than worst-case
- In practice often get  $O(b^{m/2})$  rather than  $O(b^m)$  [ $b$  =max. number of child.,  $m$  =max. depth of any node]
  - e.g., in chess go from  $b \sim 35$  to  $m \sim 6$ 
    - this permits much deeper search in the same amount of time

## Final Comments about Alpha-Beta Pruning

- Pruning does not affect final results
- Entire subtrees can be pruned
- Good move *ordering* improves effectiveness of pruning
- Repeated states are again possible
  - Store them in memory

## Example: the exact mirror image of the first example




which nodes can be pruned?



## Example: the exact mirror image of the first example

which nodes can be pruned?




Answer: **LOTS!** Because the most favorable nodes for both are explored **first** (i.e., in the diagram, are on the left-hand side).

# Example

---

which nodes can be pruned?






Answer: **NONE!** Because the most favorable nodes for both are explored **last** (i.e., in the diagram, are on the right-hand side).

# Game-tree search in the game of go

A game tree's size grows exponentially with both its depth and branching factor

- Go is huge: branching factor  $\approx 200$
- Game length  $\approx 250$  to 300 moves
- Number of paths in the game tree  $\approx 10^{525}$  to  $10^{620}$
- Much too big for a normal game trees earch
- Comparison:
  - Number of atoms in universe: about  $10^{80}$
  - Number of particles in universe: about  $10^{87}$



# Game-tree search in the game of go

- During the past couple years, go programs have gotten much better
- Main reason: Monte Carlo roll-outs
- Basic idea: do a minimax search of a randomly selected subtree
- At each node that the algorithm visits,
  - It randomly selects some of the children  
There are heuristics for deciding how many
  - Calls itself recursively on these, ignores the others

# Summary

---

- Game playing is best modeled as a search problem
- Game trees represent alternate computer/opponent moves
- Evaluation functions estimate the quality of a given board configuration for the Max player.
- Minimax is a procedure which chooses moves by assuming that the opponent will always choose the move which is best for them
- Alpha-Beta is a procedure which can prune large parts of the search tree and allow search to go deeper
- For many well-known games, computer algorithms based on heuristic search match or out-perform human world experts.

# Homework

---

Design an algorithm for tic-tac-toe

- Including minimax algorithm
- Evaluation function
- Alpha-beta pruning