


```
next_beam[n_prefix] = n_p

beam = sorted(next_beam.items(), key=lambda x: x[1], reverse=True)
beam = beam[:beam_width]

best = beam[0]
return best[0]
#####

def ctc_decode(
    log_probs: torch.LongTensor, blank: int = 0, beam_width: int = 5) -> list:
    """Decoding with CTC.

    Use beam search to approximate the maximum likelihood decoding
    from 'log_probs'. Make sure that blank tokens are removed afterwards
    and unnecessary repeated tokens are removed as well.

    Args:
        log_probs: log probabilities as defined in CTC. (shape: T x C)
        blank: The "epsilon" token that is used to represent silence.
        (integer < C, default 0)
        beam_width: The number of candidates to keep around.

    Returns:
        outputs: [y' 1, y' 2, ..., y' S], where each y'_t is between 0 and C-1.
        (shape: S (S= T))
    """
    ##### YOUR CODE GOES HERE #####
    def logsum(*nums):
        if all(n == -math.inf for n in nums):
            return -math.inf
        n_max = max(nums)
        res = math.log(sum(math.exp(n - n_max) for n in nums))
        return n_max + res

    T, C = log_probs.shape
    beam = [(tuple(), (0.0, -math.inf))]

    for t in range(T):
        next_beam = collections.defaultdict(lambda: (-math.inf, -math.inf))
        for C in range(C):
            p = log_probs[t, C]
            # p_b代表前缀不以blank结尾的概率, p_b代表以blank结尾的概率
            for prefix, (p_b, p_nb) in beam:
                if C == blank: # 如果是blank, 前缀不改变, 只有概率改变
                    n_p_b, n_p_nb = next_beam[prefix]
                    n_p_b = logsum(n_p_b, p_b + p, p_nb + p)
                    next_beam[prefix] = (n_p_b, n_p_nb)
                continue
            # 如果不是blank, 只有不以blank结尾的前缀概率改变
            end_t = prefix[-1] if prefix else None
            n_prefix = prefix + (C, )
            n_p_b, n_p_nb = next_beam[n_prefix]
            if C != end_t:
                n_p_nb = logsum(n_p_nb, p_b + p, p_nb + p)
            else:
                n_p_nb = logsum(n_p_nb, p_b + p)
                next_beam[n_prefix] = (n_p_b, n_p_nb)
            # 如果是空结尾概率, 我们这里就不变的保留
            if C == end_t:
                n_p_b, n_p_nb = next_beam[prefix]
                n_p_nb = logsum(n_p_nb, p_nb + p)
                next_beam[prefix] = (n_p_b, n_p_nb)
        beam = sorted(next_beam.items(), key=lambda x: logsum(*x[1]), reverse=True)
        beam = beam[:beam_width]

    best = beam[0]
    return best[0]
#####
```

Task 2.2 Demonstrate beam search decoding (3 points)

Print the most likely transcript for each `log_probs` and character set we provide. Use the default beam width (=5). You can load the test data by calling `get_log_probs()`. Please loop through `log_probs_batch` to get `log_probs` input to test your beam search implementation. You will output 10 likely transcripts by using the test `log_probs` in `log_probs_batch`.

Note: The most likely transcript could be gibberish.

```
In [29]: import string

def get_log_probs() -> torch.LongTensor:
    """Get minibatches to test implementation
    Returns:
        lists of log_probs
    """
    torch.manual_seed(FIX_SEED)
    np.random.seed(FIX_SEED)
    random.seed(FIX_SEED)

    T = 50 # input length
    N = 10 # batch size
    C = 27 # Class size

    # "CTC model" probabilities
    log_probs_batch = torch.randn(N, T, C).log_softmax(2).detach().requires_grad_()
    return log_probs_batch

log_probs_batch = get_log_probs() # Shape: N x T x C (10, 50, 27)

log_probs_list = get_log_probs()
char_set = list(string.ascii_lowercase) # lowercase alphabet
char_set.insert(0, "eps") # add blank as the first element
print(char_set)
```

```
In [30]: ##### YOUR CODE GOES HERE #####
N = log_probs_batch.shape[0]
for i in range(N):
    pred_seq = ctc_decode(log_probs_batch[i, :, :], beam_width = 5)
    pred_seq = [char_set[_] for _ in pred_seq]
    pred_seq = ''.join(pred_seq)
    print(pred_seq)

#####

zcitlhbajdyziostzvwchqagtctjzykrpgkosyjt
ykslbnkqzhtwtfvodqizhagldkdaizqenodvtwvrdxtqre
rzedcjuxoxmcszemnzsqvudpvvjfxubnfuwaiokauhrsabxb
dwnxlbseroudtquafeflgyvntnkatzngpyygrztkfnab
iudltveubgpathggmbmajngdfrlbuveyjaipmq
yhcnpvhwclomufngxqvmghorfywufkhatarstxtktbgppc
zqbyjngmoamfeioenkyseufcewcpdpsvbwvowqukf
vmcmchnrdwdvtwhgblpmyszfpgfaxvyszdbwiautkgtakal
khakqcoasqacxekmjbctwjmghuifymgdacsvsngwnzji
pxrkpwesuzugkqodqifkuvlcjqzqjfsadrhwjkuakfavy
```

Task 2.3 Demonstrate narrowed beam search (2 points)

Print the most likely transcript for each `log_probs` and character set we provide. Use a narrow beam (=2) and comment on any difference you see with the narrow beam as compared to using the default beam size above. Do you find the same sequences? You can load the test data by calling `get_log_probs()`. Please loop through `log_probs_batch` to get `log_probs` input to test your implementation. You will output 10 likely transcripts by using the test `log_probs` in `log_probs_batch`.

Note: The most likely transcript could be gibberish.

```
In [31]: import string

log_probs_list = get_log_probs()
char_set = list(string.ascii_lowercase) # lowercase alphabet
char_set.insert(0, "eps") # add blank as the first element
print(char_set)

['eps', 'a', 'b', 'c', 'd', 'e', 'f', 'g', 'h', 'i', 'j', 'k', 'l', 'm', 'n', 'o',
 'p', 'q', 'r', 's', 't', 'u', 'v', 'w', 'x', 'y', 'z']

In [32]: ##### YOUR CODE GOES HERE #####
N = log_probs_batch.shape[0]
for i in range(N):
    pred_seq = ctc_decode(log_probs_batch[i, :, :], beam_width = 2)
    pred_seq = [char_set[_] for _ in pred_seq]
    pred_seq = ''.join(pred_seq)
    print(pred_seq)

#####

zcitlhbajdyziostzvwchqagtctjzykrpgkosyjt
ykslbnkqzhtwtfvodqizhagldkdaizqenodvtwvrdxtqre
rzedcjuxoxmcszemnzsqvudpvvjfxubnfuwaiokauhrsabxb
dwnxlbseroudtquafeflgyvntnkatzngpyygrztkfnab
iudltveubgpathggmbmajngdfrlbuveyjaipmq
yhcnpvhwclomufngxqvmghorfywufkhatarstxtktbgppc
zqbyjngmoamfeioenkyseufcewcpdpsvbwvowqukf
vmcmchnrdwdvtwhgblpmyszfpgfaxvyszdbwiautkgtakal
nohakqcoasqacxekmjbctwjmghuifymgdacsvsngwnzji
pxrkpwesuzugkqodqifkuvlcjqzqjfsadrhwjkuakfavy
```

Answer

I find many same sequences, and I think the reason is that the probability dataset we build is too small and we use the softmax operation, so the likelihood can be very extreme. Thus the narrowed beam search can also find similar answers.

However, when it applied to real scenes, I think the narrowed beam search is more likely to lead to a poor performance, resulting a suboptimal solution.

This is the end of HW2. Great work!